

DLCV HW4 Report

tags: DLCV

Course	Student ID	Name	Date
2022 Fall NTU DLCV	R10943109	Shiuan-Yun Ding	2022-12-13


HackMD Link: <https://hackmd.io/@mirkat1206/B17X7Evvj>

Problem 1: 3D Novel View Synthesis (50%)

Metrix	Baseline (10%)	My Results
PSNR	35	35.19
SSIM	0.97	0.9742

(5%) Please explain

a. the NeRF idea in your own words



- NeRF tries to use a deep fully-connected neural network without any convolution layer to synthesize views of an object.
- NeRF takes a set of images of an object from different angle as inputs of training. For inference, it takes a 5D coordinate (x, y, z, θ, ϕ) as inputs, and outputs a single volume density and view-dependent RGB color, which can be rendered as an image of view of the object from that specific position and angle.
- I think this problem is very useful for movie/cartoon/game industry, because they need to render 3D models a lot.

b. which part of NeRF do you think is the most important


Two improvements in NeRF make it outstands the other methods:

- Positional Encoding
 - The author of NeRF states that the network that directly takes a 5D coordinate (x, y, z, θ, ϕ) as inputs cannot generate high quality results. In particular, this method fails at images with high-frequency variation in color and geometry.
 - NeRF copes with this problem by using the idea of positional encoding. It first map the original 5D coordinate (x, y, z, θ, ϕ) into a higher dimensional space, then sends the encoded coordinates into MLP.
- Hierarchical volume sampling
 - The author points out that the previous works are not efficient enough because they spend much time on sampling redundant points.
 - NeRF proposed a hierarchical volume sampling method that first optimize a "coarse" network, and based on the "coarse" networks, NeRF can produce a "fine" network that uses a more informed sampling points.
 - Compared with the previous works that use uniform sampling points, NeRF with biased, nonuniform sampling points can render a better image with high efficiency.

c. compare NeRF's pros/cons w.r.t other novel view synthesis work

- Pros

- NeRF can render an images with highest quality. The details in both geometry and apperance remain clear with NeRF, while other works produce blurred details.
- Require only 5 MB for the entire network weights.
- Cons
 - The original NeRF requires at least 12 hours for training, while LLFF only takes 10 minutes to train.



Comparison with previous works

- NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis (<https://arxiv.org/pdf/2003.08934.pdf>)
- Direct Voxel Grid Optimization: Super-fast Convergence for Radiance Fields Reconstruction (<https://arxiv.org/pdf/2111.11215.pdf>)
- NeRF神经辐射场ECCV2020 (https://blog.csdn.net/qq_40943760/article/details/125189835)

(10%) Describe the implementation details of Direct Voxel Grid Optimization (DVGO)

Inspired by NeRF, DVGO proposes a super fast convergence approach that can train a model with comparable quality with NeRF but reduce the training length from over 12 hours to less than 15 minutes.

Compared with NeRF using 5D coordinate (x, y, z, θ , ϕ), DVGO uses voxel-grid representation (x, V), where x is the queried 3D point, V is the voxel grid.

For implementation, same as NeRF, DVGO first searches the coarse geometry of a scene, then reconstructs the fine detail including view-dependent effects. The reasons that DVGO is much faster than NeRF are as the following:

- Use post-activation interpolation on voxel density to produce sharp surfaces in lower grid resolution.
- Prone the original voxel density optimization to suboptimal geometry solutions.

By using the above method, DVGO achieves ~ 45x speedup.

(15%) Evaluate the generated images and ground truth images

Three metrics

#	# of fine-tune iter	# of voxel	step size	PSNR	SSIM	LPIPS
0 (default)	20000	160 ** 3	0.5	35.1671	0.9744	0.0411
1	20000	160 ** 3	0.75	35.0645	0.9738	0.0424
2	20000	160 ** 3	0.25	35.2174	0.9745	0.0413
3	40000	160 ** 3	0.25	35.2054	0.9746	0.0407
4	20000	240 ** 3	0.5	35.3512	0.9759	0.0382
5	20000	320 ** 3	0.5	35.3325	0.9759	0.0376

Discussion

- From the experiments (0, 1, 2), we can see that step size does not make much difference.
- From the experiments (2, 3), we can see that # of fine-tune iter does not make much difference.
- From the experiments (0, 4, 5), we can see that # of voxel makes much difference.
- Even though the experiments (4, 5) have the highest performance, the training cost and storage cost are the highest as well. Therefore, I choose the experiment 0 as the final version of this homework.

Explain the meaning of three metrics

- PSNR (Peak Signal to Noise Ratio)
 - The ratio between the maximum possible power of a signal and the power of corrupting noise that affects the fidelity.
 - The bigger, the better.
- SSIM (Structural Similarity Index)
 - It measures the "structural information" similarity between two images.
 - The value is from 0 (worst) to 1 (best).
- LPIPS (Learned Perceptual Image Patch Similarity)
 - Evaluate the distance between image patches.
 - Higher means more different, while lower means more similar.
- reference:
 - 有真实参照的图像质量的客观评估指标:SSIM、PSNR和LPIPS (<https://zhuanlan.zhihu.com/p/309892873>)
 - <https://github.com/richzhang/PerceptualSimilarity>
 - Wikipedia

Problem 2: Self-Supervised Pre-training for Image Classification (50%)

	Simple Baseline (5%)	Strong Baseline (5%)	My Result
Accuracy	0.36	0.40	0.4557

(10%) Describe the implementation details

Pre-training (Backbone)

- Method : SSL Method BYOL <https://github.com/lucidrains/byol-pytorch>
- Backbone Model : Resnet50

- | Data Augmentation |
|--|
| <pre>transform=transforms.Compose([transforms.Resize(128), transforms.CenterCrop(128), transforms.ToTensor(), transforms.Normalize(mean, std),])</pre> |

- | Optimizer | Learning Rate | Batch Size | Epochs |
|-----------|---------------|------------|--------|
| Adam | 3e-4 | 64 | 47 |

Fine-tuning (Classifier)

- | Classifier Model |
|---|
| <pre>self.classifier = nn.Sequential(nn.Linear(1000, 512), nn.BatchNorm1d(512), nn.ReLU(), nn.Dropout(0.5), nn.Linear(512, 256), nn.BatchNorm1d(256), nn.ReLU(), nn.Dropout(0.5), nn.Linear(256, 65))</pre> |

- | Data Augmentation |
|---|
| <pre>transform=transforms.Compose([transforms.Resize((128, 128)), transforms.CenterCrop((128)), transforms.ColorJitter(brightness=0.5, contrast=0.5, saturation=0.5, hue=0.5), transforms.RandomHorizontalFlip(), transforms.RandomRotation((-30, 30), expand=False), transforms.ToTensor(), transforms.Normalize(mean, std),])</pre> |

- | Optimizer | Learning Rate | Batch Size | Epochs |
|-----------|---------------|------------|--------|
| Adam | default | 64 | 200 |

- Reference
 - BYOL: Bootstrap Your Own Latent: A New Approach to Self-Supervised Learning (Paper Explained) (<https://www.youtube.com/watch?v=YPfUiOMYOEI>)

(20%) Complete the table

Setting	Pre-training (Mini-ImageNet)	Fine-tuning (Office-Home dataset)	Validataion Accuracy (Office-Home dataset)
A	-	train backbone + classifier	Train: 3889/3951 (98.43%) Test: 148/406 (36.45%)
B	w/ label (by TA)	train backbone + classifier	Train: 3925/3951 (99.34%) Test: 188/406 (46.31%)
C	w/o label (by me)	train backbone + classifier	Train: 3892/3951 (98.51%) Test: 185/406 (45.57%)
D	w/ label (by TA)	fix backbone train classifier	Train: 683/3951 (17.31%) Test: 57/406 (14.04%)
E	w/o label (by me)	fix backbone train classifier	Train: 674/3951 (17.06%) Test: 55/406 (13.55%)

Discussion

- Before doing data augmentation, the fine-tuning method would easily fall into mode collapse.
- After doing data augmentation, we know that, from the difference between accuracies of train set and test set, the models A/B/C fall into over-fitting problem.
- The resulting differences between B/C and between D/E are not very obvious. Based on other students' results, C should be better than B, and E should be better than D. The reason that my result does not show this trend may be because
 - i. I did not train my backbone well enough
 - ii. Both B/C fall into over-fitting
 - iii. both D/E need more epochs for training
- However, my results have passed the baseline, I decide not to futher optimize my model.

